

Do.It.Yourself AI

RUNNING YOUR OWN AI INFERENCE SERVER

NZ Tech Expo 2026

CHAT WITH OUR AI — RIGHT NOW



[QR CODE]

Running on our own hardware. No ChatGPT. No OpenAI.

Scan now and start chatting — I'll explain what's happening under the hood while you do.

Powered by a mix of technologies spanning **38 years** — IRC (1988) meets local AI inference (2026).

WHAT IS AI?

IT DEPENDS WHO YOU ASK.

WHO	WHAT THEY SAY
 A CEO	"A productivity multiplier for my team"
 A doctor	"A tool that helps me diagnose faster"
 A researcher	"Statistical pattern matching on massive datasets"
 A DevOps engineer	"An API I can self-host and plug into my pipeline"
 A developer	"A process I can call and add into my other processes"

They're all right.

But today we're talking about running it yourself.

AI IS A 5-LAYER CAKE

— Jensen Huang, NVIDIA · blogs.nvidia.com



APPLICATIONS

IRC bots · AI agents · coding assistants · chat apps · your own tools

👉 WE BUILD HERE



MODELS

Llama 3 · Qwen · DeepSeek · Mistral · Phi · Gemma · open weights

👉 WE CHOOSE HERE



INFRASTRUCTURE

GPU server · llama.cpp · LiteLLM · OpenAI-compatible API

👉 WE RUN HERE



CHIPS

NVIDIA · AMD · Apple Silicon




ENERGY

power · cooling

"Every successful application pulls on every layer beneath it — all the way down to the power plant."

Big tech is spending **trillions** building this stack. You can join the **top three layers** for the cost of a **second-hand GPU**.

FOR A DEVELOPER, AI IS A PROCESS



HARDWARE? LESS THAN YOU THINK.



"The best hardware is what you have immediately available without spending any money"


What fits where:

SIZE	MODEL	VRAM	RUNS ON
XXS	Qwen2.5-0.5B	~300 MB	Phone, Raspberry Pi
XS	Phi-4-mini	~2 GB	Any laptop, CPU-only
M	Llama 3.1 8B	~5 GB	Gaming GPU
L	Qwen3.7 MoE	~20 GB	2× Quadro RTX 4000 👉 us
XL	Qwen3 32B	~24 GB	RTX 4090 / 3090
XXL	Qwen3.8 Max	~40 GB	Multi-GPU server

Source: [benchmark.ai](https://www.benchmark.ai) · best open weight · Aug 15, 2026

- **VRAM** is the bottleneck — fit the model in for full speed
- RAM offload works but is 10–20× slower
- CPU handles tokenization + routing + I/O

What we actually run:



- **Primary:** Qwen3.7 MoE · 22B active
- **Router:** LiteLLM — OpenAI-compatible API
- **Fallback:** Groq / DeepSeek if both busy

WHAT DOES OWNING THE INFRASTRUCTURE GIVE YOU?



PRIVACY

Your data never leaves the building

- Prompts stay on your network
- No vendor logging your queries
- GDPR / compliance friendly



COST

Pay once — not per token, forever

- Break-even in weeks vs. API
- Unlimited calls, zero surprise bills



CONTROL

You decide everything

- Swap models in seconds
- No rate limits or outages
- Fine-tune on your own data

Cloud AI charges per token — **every word, every character, every call.**

Local inference only costs **electricity.**

SCAN. CHAT. ASK US ANYTHING.



[QR CODE]

The agent you're talking to is running on our hardware — right now.

[your contact / site](#)